

# revised\_submission\_q\_learning

November 11, 2024

## 1 Q-Learning on Viral Load Data for Optimal Treatment Strategies

### 1.1 Importing Libraries

```
[1]: import pandas as pd
import numpy as np
from sklearn.preprocessing import MinMaxScaler
import random
import itertools
```

```
[2]: data = pd.read_excel('/Users/hamiddastgir/Desktop/mgddl16.xlsx')
```

```
[3]: data = data.drop(columns='Unnamed: 0')
```

```
[4]: data.head()
```

```
[4]:
```

	CASEID	VISIT		DNAME	DCLASS	NRTI	NNRTI	PI	OTHER	\
0	1000	40		Atripla	NNRTI NRTI	1	1	0	0	
1	1000	41		Atripla	NNRTI NRTI	1	1	0	0	
2	1000	42		Triumeq	OTHER	0	0	0	1	
3	1003	9	Retrovir	Epivir	NRTI NRTI	1	0	0	0	
4	1003	10	Retrovir	Epivir	NRTI NRTI	1	0	0	0	

  

	MISSING	VLOAD	...	LYMPLC	LYMPLC_INC	LPCTLC	LPCTLC_INC	LPMLC	\
0	0	20.0	...	3	0.0	31.0	0.000000	-1	
1	0	20.0	...	3	0.0	23.0	-0.258065	-1	
2	0	20.0	...	3	0.0	33.0	0.434783	-1	
3	0	8200.0	...	3	0.0	22.2	0.000000	700	
4	0	3200.0	...	3	0.0	32.2	0.450450	1200	

  

	LPMLC_INC	POLYLC	POLYLC_INC	PPCTLC	PPCTLC_INC
0	0.000000	3	0.0	58.0	0.000000
1	0.000000	3	0.0	66.0	0.137931
2	0.000000	3	0.0	56.0	-0.151515
3	0.000000	3	0.0	64.9	0.000000
4	0.714286	3	0.0	48.9	-0.246533

[5 rows x 51 columns]

## 1.2 Defining State and Action Spaces

```
[5]: state_columns = ['HGB_LC', 'MCV_LC', 'PLATLC', 'WBC_LC', 'HSRAT']
     states = data[state_columns].values

[6]: action_columns = ['NRTI', 'NNRTI', 'PI', 'OTHER']
     action = data[action_columns].values

[7]: scaler = MinMaxScaler()
     normalized_states = scaler.fit_transform(states)

[8]: action_combinations = list(itertools.product([0, 1],
     ↪repeat=len(action_columns)))

[9]: action_mapping = {combo: idx for idx, combo in enumerate(action_combinations)}
     inverse_action_mapping = {idx: combo for combo, idx in action_mapping.items()}

[10]: action_tuples = [tuple(a) for a in action]
     action_indices = [action_mapping[a] for a in action_tuples]

[11]: num_actions = len(action_combinations)
```

## 1.3 Discretizing the State Space

```
[12]: num_bins = 10
     state_bins = [np.linspace(0, 1, num_bins + 1) for _ in range(normalized_states.
     ↪shape[1])]

[13]: def discretize_state(state):
     discretized_state = []
     for i in range(len(state)):
         discretized_value = np.digitize(state[i], state_bins[i]) - 1
         discretized_state.append(discretized_value)
     return tuple(discretized_state)
```

## 1.4 Initializing the Q-Table

```
[14]: state_space = set()
     for state in normalized_states:
         state_space.add(discretize_state(state))
     state_mapping = {state: idx for idx, state in enumerate(state_space)}
     inverse_state_mapping = {idx: state for state, idx in state_mapping.items()}

     num_states = len(state_mapping)
     q_table = np.zeros((num_states, num_actions))
```

## 1.5 Setting Q-Learning Parameters

```
[15]: learning_rate_alpha = 0.1
discount_factor_gamma = 0.5
exploration_rate_epsilon = 0.1
# optional: exploration_rate_epsilon = min_epsilon + (max_epsilon -
↪ min_epsilon) * np.exp(-decay_rate * episode)
```

```
[16]: VLOAD = data['VLOAD'].values
```

## 1.6 Implementing the Q-Learning Algorithm

```
[17]: for episode in range(1000):
    for i in range(len(normalized_states) - 1):
        state = normalized_states[i]
        discretized_state = discretize_state(state)
        state_index = state_mapping[discretized_state]

        if np.random.uniform(0, 1) < exploration_rate_epsilon:
            action_index = np.random.choice(num_actions)
        else:
            action_index = np.argmax(q_table[state_index, :])

        actual_action_index = action_indices[i]

        agent_action_vector = np.array(action_combinations[action_index])
        actual_action_vector = np.
↪ array(action_combinations[actual_action_index])

        action_similarity = np.sum(agent_action_vector == actual_action_vector)
↪ / len(action_columns)

        reward = VLOAD[i] - VLOAD[i + 1]

        adjusted_reward = reward * action_similarity

        next_state = normalized_states[i + 1]
        discretized_next_state = discretize_state(next_state)
        next_state_index = state_mapping[discretized_next_state]
        next_max = np.max(q_table[next_state_index, :])

        q_table[state_index, action_index] = (1 - learning_rate_alpha) *
↪ q_table[state_index, action_index] + \
                                learning_rate_alpha *
↪ (adjusted_reward + discount_factor_gamma * next_max)
```

## 1.7 Creating DataFrames and Saving Results

```
[18]: q_table_df = pd.DataFrame(q_table, columns=[str(combo) for combo in
↳ action_combinations])
q_table_df
```

```
[18]:      (0, 0, 0, 0)  (0, 0, 0, 1)  (0, 0, 1, 0)  (0, 0, 1, 1)  (0, 1, 0, 0) \
0      69.634334    58.043931    102.175767    94.646888    9.037365
1      263.279839    277.466025    419.842723    260.095815    607.234597
2     -7734.562584   3404.775431   -8006.327947   -4153.504642   -4605.789979
3     20862.211110   24836.079175   25552.660389   22078.935838    7662.993823
4     9889.184049  -58626.240855   -9539.181734   -6093.851737  -11214.149192
..      ...
280     917.656118    1669.448479    845.843631    3114.733673    5177.531320
281    -917.453615   -379.890305   -396.433663    136.309613   -1639.619523
282   -1354.052400   -384.191501   -588.766254   -526.364593   -570.814610
283   -129.028473    47.351375   -175.747250   -28.049956    37.010977
284    335.168847    151.964422    185.943965    173.480070    189.928580

      (0, 1, 0, 1)  (0, 1, 1, 0)  (0, 1, 1, 1)  (1, 0, 0, 0)  (1, 0, 0, 1) \
0      16.543258    69.643344    47.528064    181.116367    64.071436
1      274.946304    269.090210    224.038983    250.627027    264.932540
2     -4937.870687   -4396.329168  -12397.463563  -11718.585998   -4762.859195
3      9365.134870   21752.018974   1813.115387  -12583.415776    8030.017406
4     -5968.770155   -7921.884869  -12625.374832   -5556.273060  -21053.399164
..      ...
280    1846.143125   -6347.621435   1506.972485   1185.285034   -2656.070849
281   -1100.933124   -352.703999   -628.557844   -272.813014   -421.808592
282   -341.329358   3702.616943   -389.760018  -1161.085355  -2806.834219
283    218.537955    42.008808    49.064695   2193.305316    54.927832
284    108.456918   122.524434    56.045874   136.222487   158.204773

      (1, 0, 1, 0)  (1, 0, 1, 1)  (1, 1, 0, 0)  (1, 1, 0, 1)  (1, 1, 1, 0) \
0      147.966045    35.471849    85.917438    39.954635    121.153967
1      254.199115    310.865897    290.954647    303.525054    50.074939
2     -4630.450852  -11695.482990   -5600.824425   -8951.515746  -12290.964189
3     -5006.865890   25741.650393    9174.212216  -21047.936557  -12874.564372
4    -40409.088841  -20107.165681  -15232.251940   -9040.344493  -17376.820240
..      ...
280    -450.743323   1856.697386   152.888464   -5406.024944  -36336.396487
281   -135.138036   -285.782453  -1175.216524   -426.614180   -993.860506
282   -690.940710  -1427.383014  -1889.837878  -1403.117139  -1296.746486
283     29.474240    30.202012   -71.615575   117.605229    38.172827
284    212.309424   146.584003   107.810551   140.823089   179.208479

      (1, 1, 1, 1)
0      37.033890
```

```

1      144.382188
2     -4819.174043
3     -387.307192
4    -40158.239140
...
280 -26228.901328
281  -1218.098167
282  -3523.497773
283     62.356924
284     50.254248

```

[285 rows x 16 columns]

```

[20]: state_index_to_normalized_state = {}
      for state in normalized_states:
          discretized_state = discretize_state(state)
          state_index = state_mapping[discretized_state]
          state_index_to_normalized_state[state_index] = state

      normalized_states_list = [state_index_to_normalized_state[idx] for idx in
                               ↪state_indices]

      normalized_state_df = pd.DataFrame(normalized_states_list, columns=[col +
                               ↪'_norm' for col in state_columns])

      full_df = pd.concat([state_df, normalized_state_df, q_table_df], axis=1)

      full_df.to_csv('q_table_with_og_states.csv', index=False)

```

[ ]:

[ ]: